

Lab Report

ECPE 170 – Computer Systems and Networks – Spring 2017

Name: STEVE GUERRERO

Lab Topic: PERFORMANCE MEASUREMENT (LAB #: 05)

- (1) Create a table that shows the real, user, and system times measured for the bubble and tree sort algorithms.

TIME TYPES	BUBBLE SORT TIME (100,000) elements	TREE SORT TIME (100,000) elements
REAL	34.088 seconds	0.065 seconds
USER	34.020 seconds	0.052 seconds
SYSTEM	0.008 seconds	0.012 seconds

- (2) In the sorting program, what actions take user time?

The processing time that the cpu took to complete the sorting algorithm. The disk retrieving and writing data.

- (3) In the sorting program, what actions take kernel time?

Operating system services requested by the process.

- (4) Which algorithm is the fastest? *(It's so obvious, you don't even need a timer)*

The binary tree sort is much faster than bubble sort.

- (5) Create a table that shows the total Instruction Read (IR) counts for the bubble and tree sort routines. (In the text output format, IR is the default unit used. In the GUI program, un-check the "% Relative" button to have it display absolute counts of instruction reads).

ARRAY SIZE 100000

	Bubble Sort	Tree Sort
Total IR	164,783,180,820	98,717,805

- (6) Create a table that shows the top 3 most active functions for the bubble and tree sort programs by IR count (excluding `main()`) - Sort by the "self" column if you're using `kcachegrind`, so that you see the IR counts for *just* that function, and not include the IR count for functions that it calls.

Bubble Sort (Active Functions)	Tree Sort (Active Functions)
1) <code>bubbleSort</code> (164,755,225,641)	1) <code>insert_element</code> (49 373 069)
2) <code>Random_r</code> (2,596,775)	2) <code>_int_malloc</code> (14 899 996)
3) <code>VerifySort</code> (1,900,001)	3) <code>_int_free</code> (7 400 000)

- (7) Create a table that shows, for the bubble and tree sort programs, the most CPU intensive line that is part of the most CPU intensive function. (i.e. First, take the most active function for a program. Second, find the annotated source code for that function. Third, look inside the whole function code - what is the most active line? If by some chance it happens to be another function, drill down one more level and repeat.)

Algorithm	Line #	Activity	Code
Bubble Sort	40	74 999 250 000	<code>if (array_start[j-1] > array_start[j]);</code>
Tree Sort	115	9 476 760	<code>if (element < (*sr) -> element)</code>

(8) Show the Valgrind output file for the merge sort with the intentional memory leak. Clearly highlight the line where Valgrind identified where the block was originally allocated.

```
==10085== Memcheck, a memory error detector
==10085== Copyright (C) 2002-2015, and GNU GPL'd, by Julian Seward et al.
==10085== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==10085== Command: ./sorting_program merge
==10085== Parent PID: 8817
==10085==
==10085==
==10085== HEAP SUMMARY:
==10085==   in use at exit: 400,000 bytes in 1 blocks
==10085==   total heap usage: 2 allocs, 1 frees, 401,024 bytes allocated
==10085==
==10085== 400,000 bytes in 1 blocks are definitely lost in loss record 1 of 1
==10085==    at 0x4C2FB55: calloc (in /usr/lib/valgrind/vgpreload_memcheck-amd64-linux.so)
==10085==    by 0x40085C: main (sorting.c:42)
==10085==
==10085== LEAK SUMMARY:
==10085==    definitely lost: 400,000 bytes in 1 blocks
==10085==    indirectly lost: 0 bytes in 0 blocks
==10085==    possibly lost: 0 bytes in 0 blocks
==10085==    still reachable: 0 bytes in 0 blocks
==10085==    suppressed: 0 bytes in 0 blocks
==10085==
==10085== For counts of detected and suppressed errors, rerun with: -v
==10085== ERROR SUMMARY: 1 errors from 1 contexts (suppressed: 0 from 0)
```

(9) How many bytes were leaked in the buggy program?
400,000 bytes were leaked in the buggy program.

(10) Show the Valgrind output file for the merge sort after you fixed the intentional leak.

```
==10151== Memcheck, a memory error detector
==10151== Copyright (C) 2002-2015, and GNU GPL'd, by Julian Seward et al.
==10151== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==10151== Command: ./sorting_program merge
==10151== Parent PID: 8817
==10151==
==10151==
==10151== HEAP SUMMARY:
==10151==   in use at exit: 0 bytes in 0 blocks
==10151==   total heap usage: 2 allocs, 2 frees, 401,024 bytes allocated
==10151==
==10151== All heap blocks were freed -- no leaks are possible
==10151==
==10151== For counts of detected and suppressed errors, rerun with: -v
==10151== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

(11) Show the Valgrind output file for your tree sort.

```
==10183== Memcheck, a memory error detector
==10183== Copyright (C) 2002-2015, and GNU GPL'd, by Julian Seward et al.
==10183== Using Valgrind-3.11.0 and LibVEX; rerun with -h for copyright info
==10183== Command: ./sorting_program tree
==10183== Parent PID: 8817
==10183==
==10183==
==10183== HEAP SUMMARY:
==10183==   in use at exit: 0 bytes in 0 blocks
==10183== total heap usage: 100,001 allocs, 100,001 frees, 2,401,024 bytes allocated
==10183==
==10183== All heap blocks were freed -- no leaks are possible
==10183==
==10183== For counts of detected and suppressed errors, rerun with: -v
==10183== ERROR SUMMARY: 0 errors from 0 contexts (suppressed: 0 from 0)
```

(12) How many seconds of real, user, and system time were required to complete the amplify program execution with Lenna image (Lenna_org_1024.pgm)? Document the command used to measure this.

TIME TYPES	Time to amplify Lenna_org_1024.pgm image
REAL	1.394 seconds
USER	0.856 seconds
SYSTEM	0.528 seconds

(13) Research and answer: why does real time != (user time + system time)?

Since the user time is the actual processing your hardware components do. The system or 'kernel' time does the managing of those calls to read and write to a disk or put the code into the cpu between those processes. So both get summed up to make up real time.

(14) In your report, put the output image for the Lenna image titled output<somenum>.pgm.

Output2048.pgm



(15) Excluding **main()** and everything before it, what are the top three functions run by amplify executable when you **do** count sub-functions in the total? Document the commands used to measure this. Tip: Sorting by the "Incl" (Inclusive) column in kcachegrind should be what you want.

Commands:

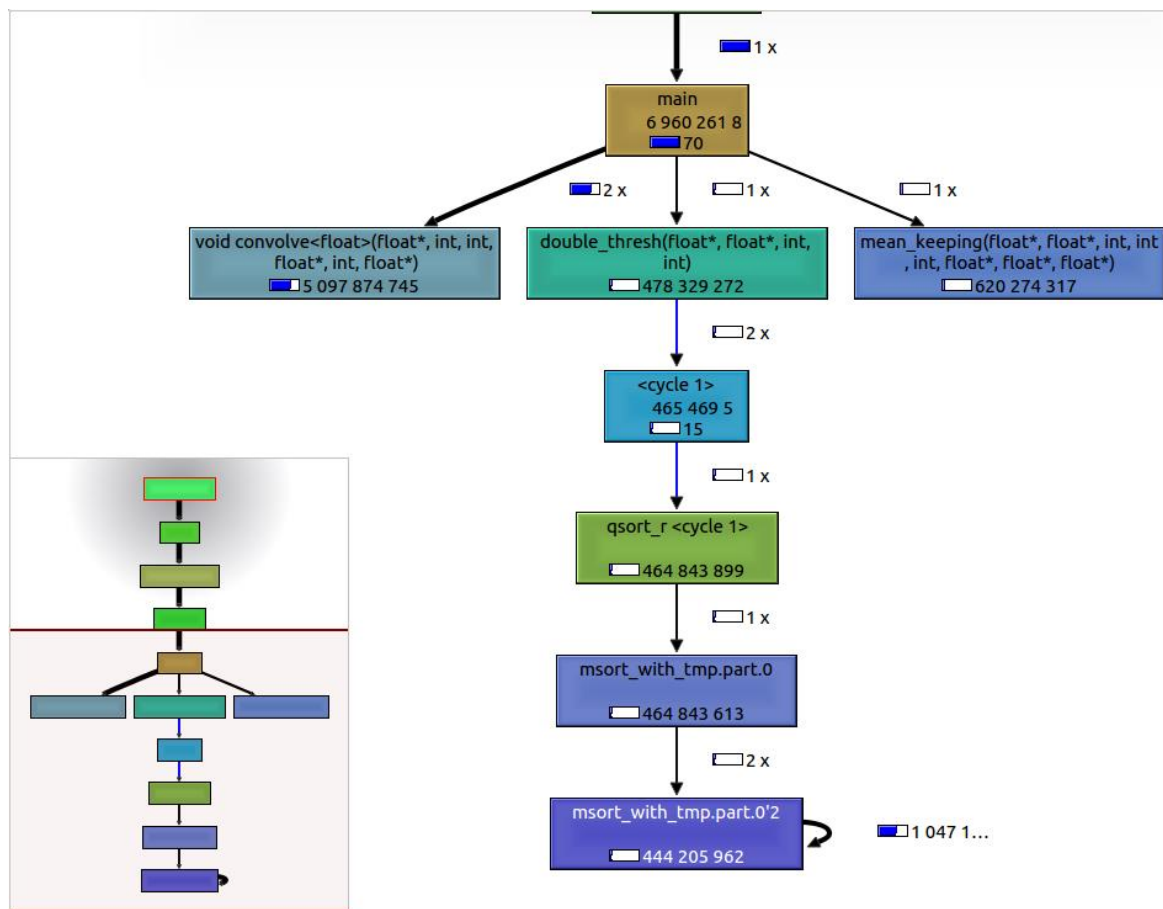
```
valgrind --tool=callgrind --dump-instr=yes --callgrind-out-file=callgrind.out ./amplify
IMAGES/Lenna_org_1024.pgm 11 1.1 2
```

```
callgrind_annotate --auto=yes --threshold=100 callgrind.out > callgrind_results.txt
```

```
kcachegrind callgrind.out &
```

Amplify Executable (Active Functions)
1) Void convolve (5 097 922 722)
2) Mean_keeping(620 280 155)
3) Double_thresh (478 333 774)

(16) Include a screen capture that shows the kcachegrind utility displaying the **callgraph** for the amplify executable, starting at main(), and going down for several levels.



(17) Which function dominates the execution time? What fraction of the total execution time does this function occupy?

Line 53: `if(((i+k) < img_height) && ((j+l) < img_width) && ((i+k)>= 0) && ((j+l)>=0))`

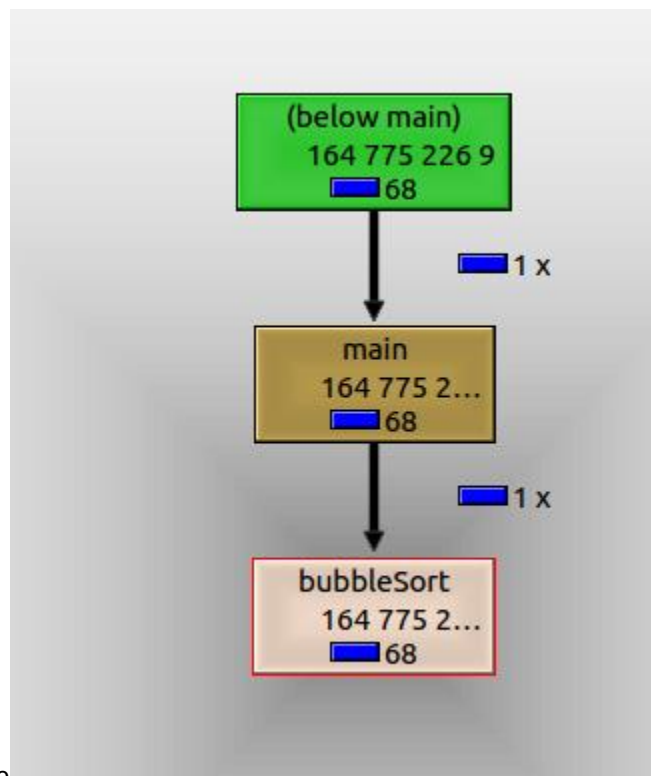
Fraction of total time:

$(2\ 348\ 521\ 224) / (6\ 962\ 390\ 673) = .3373$

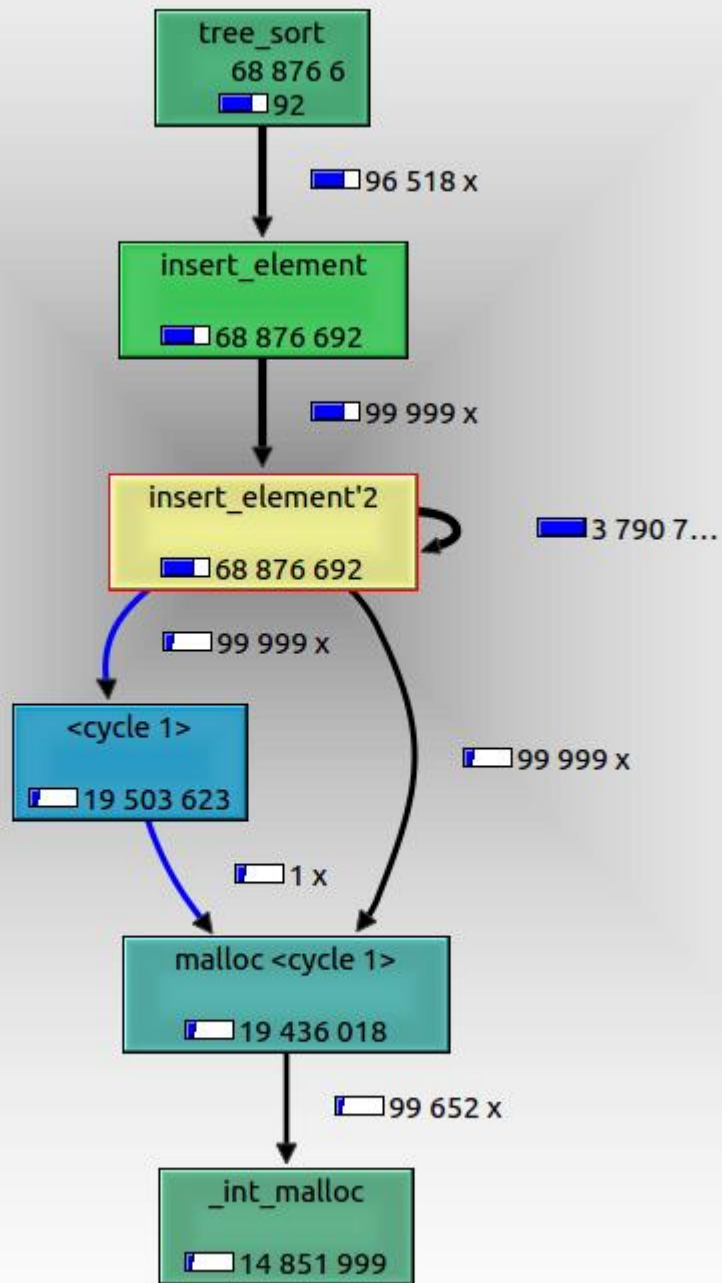
(18) Does this program show any signs of memory leaks? **Optional:** Fix it for 10 points extra credit.

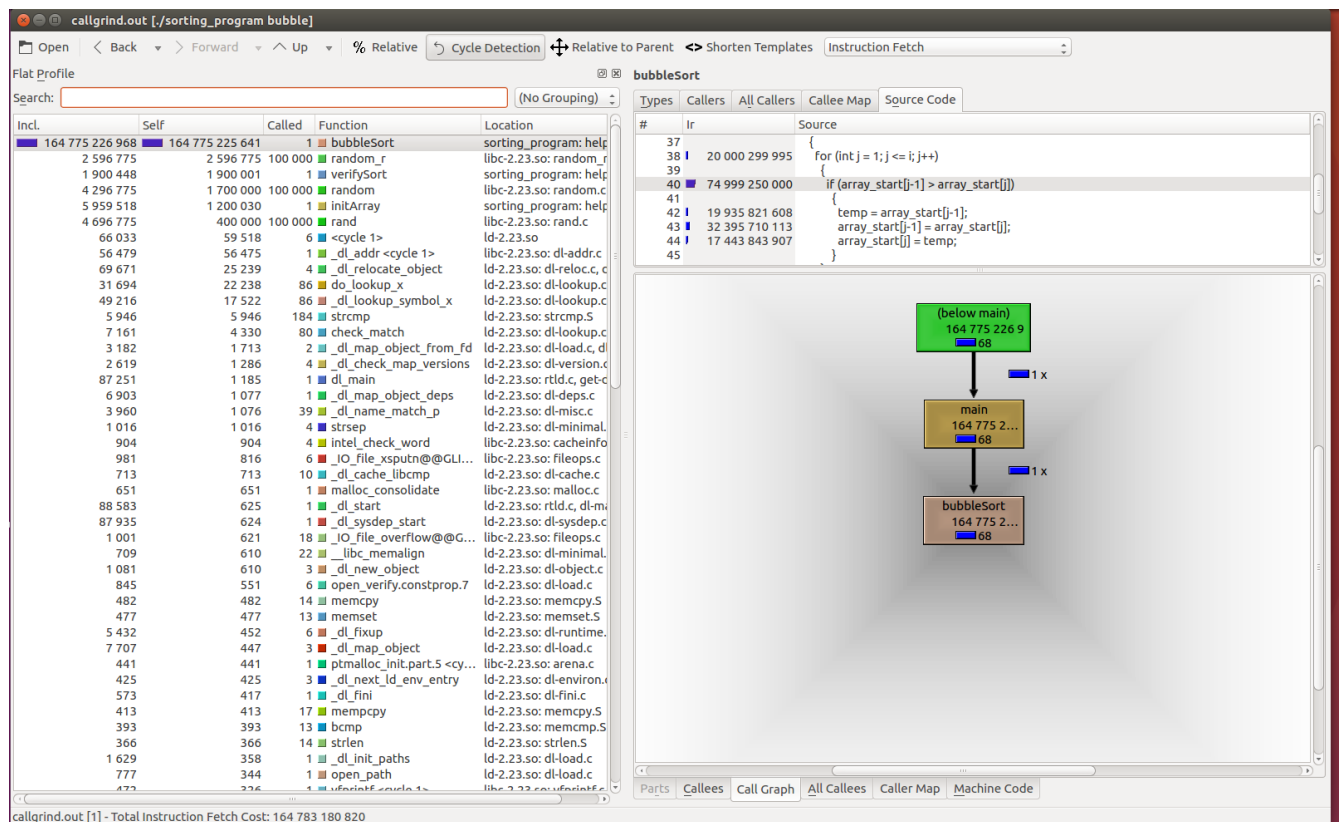
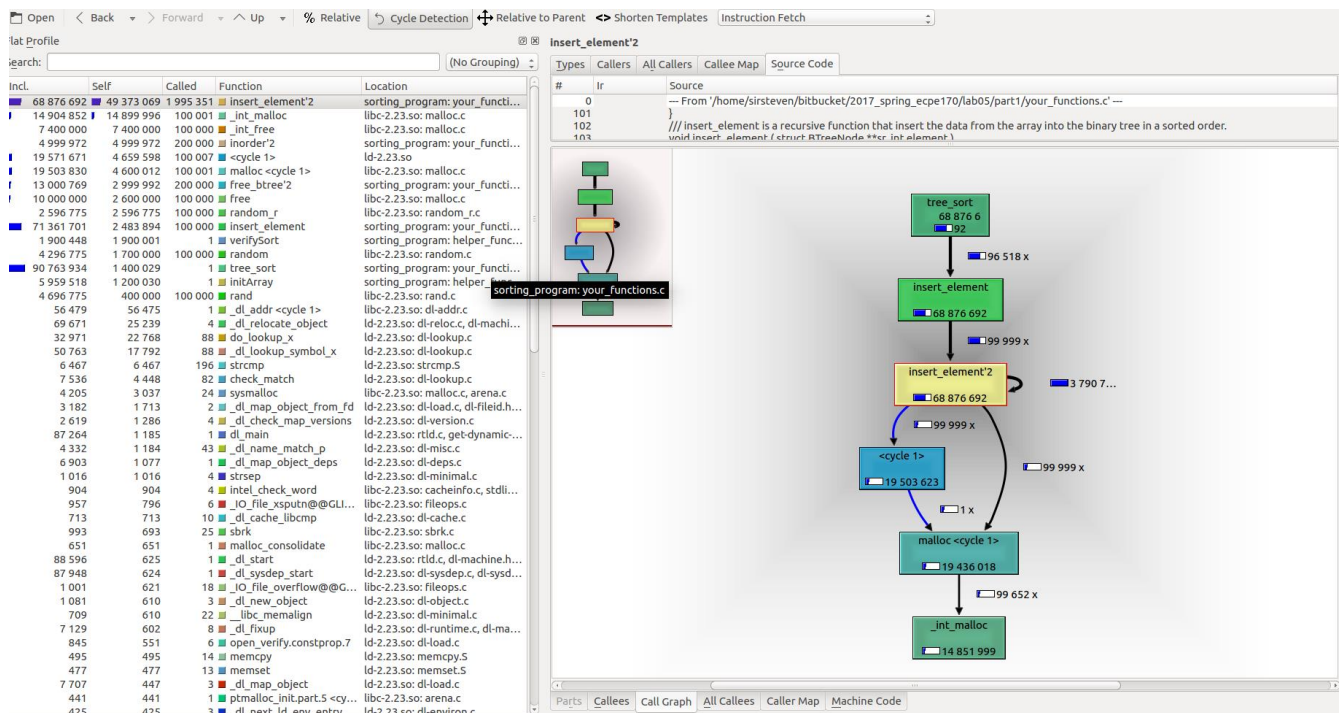
Yes, there are huge memory leaks. Memcheck shows 25,166,844 bytes in 7 blocks definitely lost.

Appendix:



BubbleSort CallTree





callgrind.out [1] - Total Instruction Fetch Cost: 164 783 180 820